

# dog\_app

November 11, 2019

## 1 Convolutional Neural Networks

### 1.1 Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

## Step 0: Import Datasets

Make sure that you've downloaded the required human and dog datasets:

**Note: if you are using the Udacity workspace, you DO NOT need to re-download these - they can be found in the /data folder as noted in the cell below.**

- Download the [dog dataset](#). Unzip the folder and place it in this project's home directory, at the location /dog\_images.
- Download the [human dataset](#). Unzip the folder and place it in the home directory, at location /lfw.

*Note: If you are using a Windows machine, you are encouraged to use [7zip](#) to extract the folder.*

In the code cell below, we save the file paths for both the human (LFW) dataset and dog dataset in the numpy arrays human\_files and dog\_files.

```
In [1]: import numpy as np
        from glob import glob

        # load filenames for human and dog images
        human_files = np.array(glob("/data/lfw/*/.*"))
        dog_files = np.array(glob("/data/dog_images/*/.*"))

        # print number of images in each dataset
        print('There are %d total human images.' % len(human_files))
        print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.

There are 8351 total dog images.

### ## Step 1: Detect Humans

In this section, we use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images.

OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the haarcascades directory. In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

```
In [20]: import cv2
        import matplotlib.pyplot as plt
        %matplotlib inline

        # extract pre-trained face detector
        face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

        # load color (BGR) image
        img = cv2.imread(human_files[0])
        # convert BGR image to grayscale
        gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

        # find faces in image
        faces = face_cascade.detectMultiScale(gray)

        # print number of faces detected in the image
        print('Number of faces detected:', len(faces))
```

```

# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 1



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

### 1.1.1 Write a Human Face Detector

We can use this procedure to write a function that returns True if a human face is detected in an image and False otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

```
In [18]: # returns "True" if face is detected in image stored at img_path
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

### 1.1.2 (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

**Answer:** (You can print out your results and/or write your percentages in this cell)

```
In [4]: from tqdm import tqdm
```

```
human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

#-#-# Do NOT modify the code above this line. #-#-#

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.
missed_humans = 0
missed_dogs = 0
for himg, dimg in tqdm(zip(human_files_short, dog_files_short)):
    missed_humans += not (face_detector(himg))
    missed_dogs += face_detector(dimg)
human_acc = (len(human_files_short) - missed_humans) / len(human_files_short)
dog_acc = (len(dog_files_short) - missed_dogs) / len(dog_files_short)
total_h_acc = (human_acc*len(human_files_short) + dog_acc*len(dog_files_short))/(len(dog
print('human_accuracy: %.2f' % human_acc)
print('dogs as human: %.2f' % (1-dog_acc))
print('Accuracy in human detection: %.3f' % total_h_acc)
```

```
100it [00:32, 3.03it/s]
```

```
human_accuracy: 0.98
dogs as human: 0.17
```

Accuracy in human detection: 0.905

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on human\_files\_short and dog\_files\_short.

```
In [5]: ### (Optional)
        ### TODO: Test performance of anotherface detection algorithm.
        ### Feel free to use as many code cells as needed.
        def check_detection_acc(face_cascade):
            human_files_short = human_files[:100]
            dog_files_short = dog_files[:100]
            missed_humans = 0
            missed_dogs = 0
            for himg, dimg in tqdm(zip(human_files_short, dog_files_short)):
                missed_humans += not (face_detection(himg, face_cascade))
                missed_dogs += face_detection(dimg, face_cascade)
            human_acc = (len(human_files_short) - missed_humans) / len(human_files_short)
            dog_acc = (len(dog_files_short) - missed_dogs) / len(dog_files_short)
            total_h_acc = (human_acc * len(human_files_short) + dog_acc * len(dog_files_short)) / (len(human_files_short) + len(dog_files_short))
            return human_acc, dog_acc, total_h_acc

        def face_detection(img_path, face_cascade):
            img = cv2.imread(img_path)
            gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
            faces = face_cascade.detectMultiScale(gray)
            return len(faces) > 0

In [6]: ### (Optional)
        ### TODO: Test performance of anotherface detection algorithm.
        ### Feel free to use as many code cells as needed.
        face_cascade_ = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')
        h_a, d_a, th_a = check_detection_acc(face_cascade_)
        print('human accuracy: %.2f' % h_a)
        print('dogs as human: %.2f' % (1-d_a))
        print('Accuracy in human detection: %.3f' % th_a)

100it [01:28, 3.05it/s]

human_accuracy: 0.98
dogs as human: 0.17
Accuracy in human detection: 0.905
```

```

In [ ]: ### (Optional)
        ### TODO: Test performance of anotherface detection algorithm.
        ### Feel free to use as many code cells as needed.
        face_cascade_ = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_default.xml')
        h_a, d_a, th_a = check_detection_acc(face_cascade_)
        print('human_accuracy: %.2f' % h_a)
        print('dogs as human: %.2f' % (1-d_a))
        print('Accuracy in human detection: %.3f' % th_a)

In [ ]: ### (Optional)
        ### TODO: Test performance of anotherface detection algorithm.
        ### Feel free to use as many code cells as needed.
        face_cascade_ = cv2.CascadeClassifier('haarcascades/haarcascade_smile.xml')
        h_a, d_a, th_a = check_detection_acc(face_cascade_)
        print('human_accuracy: %.2f' % h_a)
        print('dogs as human: %.2f' % (1-d_a))
        print('Accuracy in human detection: %.3f' % th_a)

```

---

## ## Step 2: Detect Dogs

In this section, we use a [pre-trained model](#) to detect dogs in images.

### 1.1.3 Obtain Pre-trained VGG-16 Model

The code cell below downloads the VGG-16 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#).

```

In [1]: import torch
        import torchvision.models as models

        # define VGG16 model
        VGG16 = models.vgg16(pretrained=True)

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()

        # move model to GPU if CUDA is available
        if use_cuda:
            VGG16 = VGG16.cuda()

```

```

Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg
100%|| 553433881/553433881 [00:05<00:00, 103950416.25it/s]

```

Given an image, this pre-trained VGG-16 model returns a prediction (derived from the 1000 possible categories in ImageNet) for the object that is contained in the image.

### 1.1.4 (IMPLEMENTATION) Making Predictions with a Pre-trained Model

In the next code cell, you will write a function that accepts a path to an image (such as 'dogImages/train/001.Affenpinscher/Affenpinscher\_00001.jpg') as input and returns the index corresponding to the ImageNet class that is predicted by the pre-trained VGG-16 model. The output should always be an integer between 0 and 999, inclusive.

Before writing the function, make sure that you take the time to learn how to appropriately pre-process tensors for pre-trained models in the [PyTorch documentation](#).

```
In [16]: from PIL import Image
import torchvision.transforms as transforms

def VGG16_predict(img_path, img_size, model):
    """
    Use pre-trained VGG-16 model to obtain index corresponding to
    predicted ImageNet class for image at specified path

    Args:
        img_path: path to an image

    Returns:
        Index corresponding to VGG-16 model's prediction
    """
    ## TODO: Complete the function.
    ## Load and pre-process an image from the given img_path
    img = Image.open(img_path)
    in_transform = transforms.Compose([
        transforms.CenterCrop(img_size),
        transforms.ToTensor(),
        # transforms.Normalize()
    ])
    image = in_transform(img).unsqueeze(0)
    # print(image.shape)
    ## Return the *index* of the predicted class for that image
    model.eval()
    _, pred_idx = torch.max(model(image), 1)
    return pred_idx # predicted class index
```

### 1.1.5 (IMPLEMENTATION) Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained VGG-16 model, we need only check if the pre-trained model predicts an index between 151 and 268 (inclusive).

Use these ideas to complete the dog\_detector function below, which returns True if a dog is detected in an image (and False if not).

```
In [2]: ### returns "True" if a dog is detected in the image stored at img_path
```

```
def dog_detector(img_path, img_size, model = VGG16):
    ## TODO: Complete the function.
    output = VGG16_predict(img_path, img_size, model)
    return output in range(151, 268)
```

### 1.1.6 (IMPLEMENTATION) Assess the Dog Detector

**Question 2:** Use the code cell below to test the performance of your dog\_detector function.

- What percentage of the images in human\_files\_short have a detected dog?
- What percentage of the images in dog\_files\_short have a detected dog?

**Answer:**

```
In [18]: ### TODO: Test the performance of the dog_detector function
         ### on the images in human_files_short and dog_files_short.

human_files_short = human_files[:100]
dog_files_short = dog_files[:100]

human_as_dog = 0
dog_wclassified = 0

for img_path in tqdm(human_files_short):
    human_as_dog += dog_detector(img_path, 224)

for img_path in tqdm(dog_files_short):
    dog_wclassified += dog_detector(img_path, 224)
h_accuracy = (len(human_files_short)-human_as_dog)/len(human_files_short)
d_accuracy = dog_wclassified / len(dog_files_short)
print("human misclassified as dogs: %.2f %" %(1-h_accuracy))
print("dogs well-classified as dogs: %.2f %" %d_accuracy)

0%|          | 0/100 [00:00<?, ?it/s]
1%|          | 1/100 [00:01<01:45, 1.06s/it]
2%|          | 2/100 [00:01<01:38, 1.00s/it]
3%|          | 3/100 [00:02<01:33, 1.04it/s]
4%|          | 4/100 [00:03<01:29, 1.07it/s]
5%|          | 5/100 [00:04<01:27, 1.09it/s]
100%|| 100/100 [01:27<00:00, 1.15it/s]
100%|| 100/100 [01:28<00:00, 1.11it/s]

human misclassified as dogs: 0.00 %
dogs well-classified as dogs: 0.59 %
```

We suggest VGG-16 as a potential network to detect dog images in your algorithm, but you are free to explore other pre-trained networks (such as [Inception-v3](#), [ResNet-50](#), etc). Please use



the code cell below to test other pre-trained PyTorch models. If you decide to pursue this *optional* task, report performance on `human_files_short` and `dog_files_short`.

```
In [6]: ### (Optional)  
### TODO: Report the performance of another pre-trained network.  
### Feel free to use as many code cells as needed.
```

```
# define Inception V3 model  
IncV3 = models.inception_v3(pretrained=True)
```

```
# check if CUDA is available  
use_cuda = torch.cuda.is_available()
```

```
# move model to GPU if CUDA is available  
if use_cuda:  
    IncV3 = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/inception_v3_google-1a9a5a14.pth" to /root/.torch/models/100%|| 108857766/108857766 [00:01<00:00, 99235441.35it/s]
```

```
In [7]: # ResNet-50  
# define Resnet50 model  
RNET = models.resnet50(pretrained=True)
```

```
# check if CUDA is available  
use_cuda = torch.cuda.is_available()
```

```
# move model to GPU if CUDA is available  
if use_cuda:  
    RNET = VGG16.cuda()
```

```
Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/100%|| 102502400/102502400 [00:01<00:00, 87675184.45it/s]
```

```
In [9]: def check_detector_performance(input_img_size = 224, model = VGG16, detector = dog_detector):  
    human_files_short = human_files[:100]  
    dog_files_short = dog_files[:100]
```

```
    human_as_dog = 0  
    dog_wclassified = 0
```

```
    for img_path in tqdm(human_files_short):  
        human_as_dog += detector(img_path, input_img_size, model)
```

```
    for img_path in tqdm(dog_files_short):  
        dog_wclassified += detector(img_path, input_img_size, model)
```

```

        h_accuracy = (len(human_files_short)-human_as_dog)/len(human_files_short)
        d_accuracy = dog_wclassified / len(dog_files_short)
        return h_accuracy, d_accuracy

In [23]: h_acc, d_acc = check_detector_performance(224, VGG16, dog_detector)
        print("human misclassified as dogs: %.2f" %h_acc)
        print("dogs well-classified as dogs: %.2f" %d_acc)

100%|| 100/100 [01:31<00:00, 1.10it/s]
100%|| 100/100 [01:21<00:00, 1.19it/s]

human misclassified as dogs: 1.00
dogs well-classified as dogs: 0.59


In [24]: h_acc, d_acc = check_detector_performance(448, IncV3, dog_detector)
        print("human misclassified as dogs: %.2f" %h_acc)
        print("dogs well-classified as dogs: %.2f" %d_acc)

100%|| 100/100 [01:41<00:00, 1.07it/s]
100%|| 100/100 [01:33<00:00, 1.14it/s]

human misclassified as dogs: 1.00
dogs well-classified as dogs: 0.96


In [16]: h_acc, d_acc = check_detector_performance(224, RNET, dog_detector)
        print("human misclassified as dogs: %.2f" %h_acc)
        print("dogs well-classified as dogs: %.2f" %d_acc)

100%|| 100/100 [00:29<00:00, 3.41it/s]
100%|| 100/100 [00:30<00:00, 3.39it/s]

human misclassified as dogs: 1.00
dogs well-classified as dogs: 0.75

```

---

### ## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 10%. In Step 4 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have trouble distinguishing between a Brittany and a Welsh Springer Spaniel.

|          |                        |
|----------|------------------------|
| Brittany | Welsh Springer Spaniel |
|----------|------------------------|

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

|                        |                        |
|------------------------|------------------------|
| Curly-Coated Retriever | American Water Spaniel |
|------------------------|------------------------|

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

|                 |                    |
|-----------------|--------------------|
| Yellow Labrador | Chocolate Labrador |
|-----------------|--------------------|

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

#### 1.1.7 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at `dog_images/train`, `dog_images/valid`, and `dog_images/test`, respectively). You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [19]: import os
         from torchvision import datasets, transforms
         from PIL import ImageFile
         ImageFile.LOAD_TRUNCATED_IMAGES = True

         ### TODO: Write data loaders for training, validation, and test sets
```

```

## Specify appropriate transforms, and batch_sizes
def load_data_in_DataLoaders(data_dir, tr_data_transform = transforms.ToTensor(), vt_da
    # Assemble dir paths
    # train_dir = os.path.join(data_dir, 'train')
    # valid_dir = os.path.join(data_dir, 'valid')
    # test_dir = os.path.join(data_dir, 'test')
    net_stage = ['train', 'valid', 'test']
    # Load Images dataset using ImageFolder
    data_scratch = { item: datasets.ImageFolder(os.path.join(data_dir, item), transform
                    for item in net_stage}
    # train_data = datasets.ImageFolder(train_dir, transform=vt_data_transform)
    # valid_data = datasets.ImageFolder(valid_dir, transform=vt_data_transform)
    # test_data = datasets.ImageFolder(test_dir, transform=vt_data_transform)
    # Print out info
    for stage in net_stage:
        print('Num images in', stage, ': ', len(data_scratch[stage]) )
    # print('Num training images: ', len(train_data))
    # print('Num validation images: ', len(valid_data))
    # print('Num test images: ', len(test_data))
    # Set data Loaders
    # define dataloader parameters
    batch_size = 20
    num_workers=0

    # prepare data loaders
    # train_loader = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
    # valid_loader = torch.utils.data.DataLoader(valid_data, batch_size=batch_size,
    loaders_scratch = {item : torch.utils.data.DataLoader(data_scratch[item], batch_size=batch_size,
                                                            num_workers=num_workers, shuffle=True) for item in net_stage}

    return loaders_scratch

# Define transforms on the data
tr_data_transform = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.RandomHorizontalFlip(0.05),
    transforms.RandomRotation(30),
    transforms.RandomAffine(10),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std=[0.229, 0.224, 0.225])
])

tr_crop_data_transform = transforms.Compose([
    transforms.CenterCrop(224),
    transforms.RandomHorizontalFlip(0.05),
    transforms.RandomRotation(30),
    transforms.RandomAffine(10),
    transforms.ToTensor(),

```

```

        transforms.Normalize(mean=[0.485, 0.456, 0.406],
                               std=[0.229, 0.224, 0.225])
    ])
    vt_data_transform = transforms.Compose([
        transforms.Resize((224,224)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406],
                               std=[0.229, 0.224, 0.225])
    ])
    vt_crop_data_transform = transforms.Compose([
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406],
                               std=[0.229, 0.224, 0.225])
    ])

    res_loaders = load_data_in_DataLoaders('/data/dog_images', vt_data_transform, vt_data_t
    cr_loaders = load_data_in_DataLoaders('/data/dog_images', vt_crop_data_transform, vt_cr

    resaug_loaders = load_data_in_DataLoaders('/data/dog_images', tr_data_transform, vt_da
    craug_loaders = load_data_in_DataLoaders('/data/dog_images', tr_crop_data_transform, vt

Num images in train : 6680
Num images in valid : 835
Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836

In [9]: max_label = 0
        labels_num = np.zeros(133)
        for images, labels in res_loaders['train']:
            for label in labels:
                labels_num[label] += 1
                if max_label < label:
                    max_label = label

In [10]: print(max_label)
          print(labels_num)
          len(labels_num)

tensor(132)
[ 64.  58.  52.  63.  77.  64.  50.  66.  34.  50.  66.  66.  46.  69.  73.

```

```

59. 62. 50. 48. 62. 64. 47. 65. 62. 37. 41. 64. 35. 74. 52.
56. 65. 45. 64. 53. 65. 50. 57. 69. 53. 69. 63. 50. 64. 53.
67. 54. 54. 50. 50. 62. 49. 47. 57. 50. 65. 71. 50. 47. 60.
61. 53. 53. 39. 42. 33. 34. 63. 51. 47. 62. 48. 42. 41. 44.
64. 43. 40. 59. 46. 56. 61. 46. 50. 37. 53. 66. 51. 53. 58.
57. 44. 35. 44. 49. 43. 50. 46. 42. 34. 48. 29. 58. 42. 31.
50. 46. 26. 45. 33. 44. 54. 39. 35. 63. 30. 48. 53. 31. 39.
28. 32. 44. 50. 34. 30. 41. 30. 48. 44. 30. 26. 30.]

```

Out[10]: 133

**Question 3:** Describe your chosen procedure for preprocessing the data. - How does your code resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

**Answer:**

After looking into the "dog\_images" folder we can make the following statements: - There are ~8k images. Which should be initially enough to train and test a pre-trained network - Most breeds are represented by 5-8 images in the test dataset; whereas there are 50 images for training each breed - Images have different zoom but most of them contain the dog in the center of the picture without rotation (dog feet are normally at the bottom) - Full-body pictures are centered too but the head of the dog lies on one side. A crop could erase relevant information related to head shape.

The following preprocessing has been applied to enhance detection results: - Image has to be resized to 224 px to feed into VGG16 (Inception uses 448). This resizing can be done either by cropping (Which can result in erasing head or other part of the dog relevant for detection) or by resizing (Which applies perspective perturbations to the picture). As both solutions can hamper proper detection, I encourage testing both. The cell above contains loaders for both transformations. Note that crop in training is randomly performed in an attempt to reduce centerCrop effect during test and validation. - Considering a real-world future application, rotation and perspective invariance should be provided by data augmentation. In such a manner rotation, horizontal flip and affine transformation is applied on training data. To monitor the benefit of this fact, loaders without augmentation are also kept. For the test data provided, results should be really similar between augmentation and without as most test images have similar orientation and perspective than training images. - Normalization can help convergence by reducing pixel-value-variance range between images. Normalization mean and std values taken from: <https://github.com/pytorch/examples/blob/97304e232807082c2e7b54c597615dc0ad8f6173/imagenet/main.py#L198>

### 1.1.8 (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. Use the template in the code cell below.

```

In [5]: import torch.nn as nn
        import torch.nn.functional as F
        from collections import OrderedDict
        from tqdm import tqdm

```

```

# define the CNN architecture
class Net(nn.Module):
    ### TODO: choose an architecture, and complete the class
    def __init__(self, dog_breeds):
        super(Net, self).__init__()

        ## Define layers of a CNN as Ordered Dicts
        self.features = nn.Sequential(OrderedDict([
            # convolutional layer (gets 224x224x3 image tensor)
            ('conv1', nn.Conv2d(3, 16, 3, padding=1)),
            ('ReLU1', nn.ReLU(True)),
            ('maxpool1', nn.MaxPool2d(2,2)),
            # convolutional layer (gets 128x128x16 tensor)
            ('conv2', nn.Conv2d(16, 32, 3, padding=1)),
            ('ReLU2', nn.ReLU(True)),
            ('maxpool2', nn.MaxPool2d(2,2)),
            # convolutional layer (gets 64x64x32 tensor)
            ('conv3', nn.Conv2d(32, 64, 3, padding=1)),
            ('ReLU3', nn.ReLU(True)),
            ('maxpool3', nn.MaxPool2d(2,2)),
            # convolutional layer (gets 32x32x64 tensor)
            ('conv4', nn.Conv2d(64, 128, 3, padding=1)),
            ('ReLU4', nn.ReLU(True)),
            ('maxpool4', nn.MaxPool2d(2,2)),
            # convolutional layer (gets 16x16x128 tensor)
            ('conv5', nn.Conv2d(128, 256, 3, padding=1)),
            ('ReLU5', nn.ReLU(True)),
            ('maxpool5', nn.MaxPool2d(2,2)),
            # convolutional layer (gets 8 x 8 x 256)
            ('conv6', nn.Conv2d(256, 512, 3, padding=1)),
            ('ReLU6', nn.ReLU(True)),
            ('maxpool6', nn.MaxPool2d(2,2))
        ]))

        #outputs 256*8*8
        self.classifier = nn.Sequential(OrderedDict([
            # linear layer (512*4*4=8192)
            ('dropout1', nn.Dropout(0.5)),
            ('fc1', nn.Linear(512*3*3, 1000)),
            ('ReLU1', nn.ReLU(True)),
            ('dropout2', nn.Dropout(0.5)),
            # linear layer (6000 -> 1000)
            ('fc2', nn.Linear(1000, dog_breeds))
        ]))

    def forward(self, x):
        ## Define forward behavior
        x = self.features(x)

```

```

        #print(x.shape)
        x = x.view(x.shape[0],-1)
        x = self.classifier(x)
        return x

###-# You so NOT have to modify the code below this line. #-#-#

# instantiate the CNN
dog_breeds = 133
model_scratch = Net(dog_breeds)

# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

```

**Question 4:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

**Answer:**

I've chosen a kernel 3x3 as it one of the most encouraged in state of the art when working with limited resolution. Following previous examples, the first layer uses 16 filters (depth). Then, by using a maxpool layer the (X,Y) is reduced by half while the depth is doubled, promoting feature detection. I've decided to using 6 conv layer to reduce the image resolution down to 4x4 in the last layer. The networks works with 224 x 224 px so as to have similar resolution than vgg16, the network against our scrath-net competes with.

Then, we have the classifier part using fully-connected layers with ReLU activation. There are also 2 dropout layers to reduce overfitting. I have decided to apply three fully-connected layers to improve interpolation and classification as the classifier outputs 4608 features to the classifier.

### 1.1.9 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and the optimizer as `optimizer_scratch` below.

```

In [30]: import torch.optim as optim

        ### TODO: select loss function
        criterion_scratch = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_scratch = optim.SGD(model_scratch.parameters(), lr=0.03)

```

### 1.1.10 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_scratch.pt'`.

```

In [8]: def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
        """returns trained model"""
        # initialize tracker for minimum validation loss

```



```

valid_loss_min = np.Inf

for epoch in range(1, n_epochs+1):
    # initialize variables to monitor training and validation loss
    train_loss = 0.0
    valid_loss = 0.0

    #####
    # train the model #
    #####
    model.train()
    for batch_idx, (data, target) in enumerate(loaders['train']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## find the loss and update the model parameters accordingly
        # clear the gradients of all optimized variables
        #print(target)
        #print(target.shape)
        optimizer.zero_grad()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the batch loss
        loss = criterion(output, target)
        # backward pass: compute gradient of the loss with respect to model parameters
        loss.backward()
        # perform a single optimization step (parameter update)
        optimizer.step()
        # update training loss
        #train_loss += loss.item()*data.size(0)
        ## record the average training loss, using something like
        train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data - train_loss))

    #####
    # validate the model #
    #####
    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['valid']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        ## update the average validation loss
        output = model(data)
        loss = criterion(output, target)
        valid_loss = valid_loss + ((1 / (batch_idx + 1)) * (loss.data - valid_loss))

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(

```

```

        epoch,
        train_loss,
        valid_loss
    ))

    ## TODO: save the model if validation loss has decreased
    if valid_loss < valid_loss_min:
        print('Validation loss decreased. Saving model ...')
        torch.save(model.state_dict(), save_path)
        torch.save
        valid_loss_min = valid_loss

    # return trained model
    return model

In [8]: # set loaders to use
        loaders_scratch = res_loaders

        # train the model
        model_scratch = train(10, loaders_scratch, model_scratch, optimizer_scratch,
                               criterion_scratch, use_cuda, 'model_scratch.pt')

        # load the model that got the best validation accuracy
        model_scratch.load_state_dict(torch.load('model_scratch.pt'))

0%|          | 0/10 [00:00<?, ?it/s]

Epoch: 1          Training Loss: 4.889492          Validation Loss: 4.888226
Validation loss decreased. Saving model ...

20%|         | 2/10 [03:20<13:51, 103.99s/it]

Epoch: 2          Training Loss: 4.887559          Validation Loss: 4.886303
Validation loss decreased. Saving model ...

30%|         | 3/10 [04:51<11:40, 100.07s/it]

Epoch: 3          Training Loss: 4.885317          Validation Loss: 4.883965
Validation loss decreased. Saving model ...

40%|         | 4/10 [06:22<09:43, 97.30s/it]

Epoch: 4          Training Loss: 4.882019          Validation Loss: 4.879751
Validation loss decreased. Saving model ...

50%|         | 5/10 [07:52<07:56, 95.21s/it]

```

```
Epoch: 5          Training Loss: 4.873405          Validation Loss: 4.865295
Validation loss decreased.  Saving model ...
```

```
60%|      | 6/10 [09:22<06:14, 93.75s/it]
```

```
Epoch: 6          Training Loss: 4.865494          Validation Loss: 4.858145
Validation loss decreased.  Saving model ...
```

```
70%|     | 7/10 [10:53<04:38, 92.71s/it]
```

```
Epoch: 7          Training Loss: 4.857098          Validation Loss: 4.848002
Validation loss decreased.  Saving model ...
```

```
80%|    | 8/10 [12:23<03:03, 91.89s/it]
```

```
Epoch: 8          Training Loss: 4.841722          Validation Loss: 4.821765
Validation loss decreased.  Saving model ...
```

```
90%|   | 9/10 [13:53<01:31, 91.33s/it]
```

```
Epoch: 9          Training Loss: 4.790226          Validation Loss: 4.748716
Validation loss decreased.  Saving model ...
```

```
100%|| 10/10 [15:23<00:00, 90.99s/it]
```

```
Epoch: 10         Training Loss: 4.734379          Validation Loss: 4.696525
Validation loss decreased.  Saving model ...
```

### 1.1.11 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 10%.

```
In [9]: def test(loaders, model, criterion, use_cuda):

        # monitor test loss and accuracy
        test_loss = 0.
        correct = 0.
        total = 0.

        model.eval()
```

```

for batch_idx, (data, target) in enumerate(loaders['test']):
    # move to GPU
    if use_cuda:
        data, target = data.cuda(), target.cuda()
    # forward pass: compute predicted outputs by passing inputs to the model
    output = model(data)
    # calculate the loss
    loss = criterion(output, target)
    # update average test loss
    test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
    # convert output probabilities to predicted class
    pred = output.data.max(1, keepdim=True)[1]
    # compare predictions to true label
    correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
    total += data.size(0)

print('Test Loss: {:.6f}\n'.format(test_loss))

print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
    100. * correct / total, correct, total))

# call test function
test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

Test Loss: 4.687901

Test Accuracy: 2% (19/836)

```

In [10]: def evaluate_model(loaders_scratch, model_scratch, optimizer_scratch, criterion_scratch,
                           use_cuda, epochs = 10, save_path = 'model_scratch.pt'):

```

```

    # train the model
    model_scratch = train(epochs, loaders_scratch, model_scratch, optimizer_scratch,
                          criterion_scratch, use_cuda, save_path)

    # load the model that got the best validation accuracy
    model_scratch.load_state_dict(torch.load(save_path))

    # call test function
    test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)

```

```

In [12]: evaluate_model(res_loaders, model_scratch, optimizer_scratch, criterion_scratch, use_cuda,

```

10%| | 1/10 [01:30<13:34, 90.50s/it]

Epoch: 1            Training Loss: 4.684037            Validation Loss: 4.612476  
Validation loss decreased. Saving model ...

20%|            | 2/10 [03:01<12:04, 90.62s/it]

Epoch: 2            Training Loss: 4.573919            Validation Loss: 4.552939  
Validation loss decreased. Saving model ...

30%|            | 3/10 [04:32<10:34, 90.65s/it]

Epoch: 3            Training Loss: 4.505886            Validation Loss: 4.504832  
Validation loss decreased. Saving model ...

40%|            | 4/10 [06:03<09:04, 90.73s/it]

Epoch: 4            Training Loss: 4.442260            Validation Loss: 4.523919

50%|            | 5/10 [07:33<07:33, 90.77s/it]

Epoch: 5            Training Loss: 4.374512            Validation Loss: 4.457346  
Validation loss decreased. Saving model ...

60%|            | 6/10 [09:05<06:03, 90.87s/it]

Epoch: 6            Training Loss: 4.303151            Validation Loss: 4.370283  
Validation loss decreased. Saving model ...

70%|            | 7/10 [10:35<04:32, 90.82s/it]

Epoch: 7            Training Loss: 4.209877            Validation Loss: 4.290996  
Validation loss decreased. Saving model ...

80%|            | 8/10 [12:05<03:01, 90.62s/it]

Epoch: 8            Training Loss: 4.106763            Validation Loss: 4.259809  
Validation loss decreased. Saving model ...

90%|            | 9/10 [13:35<01:30, 90.43s/it]

Epoch: 9            Training Loss: 4.003677            Validation Loss: 4.200664  
Validation loss decreased. Saving model ...

100%|| 10/10 [15:06<00:00, 90.44s/it]

Epoch: 10           Training Loss: 3.894820            Validation Loss: 4.167188  
Validation loss decreased. Saving model ...

Test Loss: 4.157565

Test Accuracy: 8% (68/836)

In [13]: evaluate\_model(cr\_loaders, model\_scratch, optimizer\_scratch, criterion\_scratch, use\_cuda)

10%| | 1/10 [01:09<10:22, 69.17s/it]

Epoch: 1 Training Loss: 4.358610 Validation Loss: 4.307395  
Validation loss decreased. Saving model ...

20%| | 2/10 [02:18<09:13, 69.22s/it]

Epoch: 2 Training Loss: 4.197330 Validation Loss: 4.268797  
Validation loss decreased. Saving model ...

30%| | 3/10 [03:27<08:04, 69.20s/it]

Epoch: 3 Training Loss: 4.069749 Validation Loss: 4.222246  
Validation loss decreased. Saving model ...

40%| | 4/10 [04:36<06:54, 69.11s/it]

Epoch: 4 Training Loss: 3.958261 Validation Loss: 4.156859  
Validation loss decreased. Saving model ...

50%| | 5/10 [05:45<05:44, 68.93s/it]

Epoch: 5 Training Loss: 3.822164 Validation Loss: 4.238466

60%| | 6/10 [06:54<04:35, 68.97s/it]

Epoch: 6 Training Loss: 3.682541 Validation Loss: 4.232671

70%| | 7/10 [08:02<03:26, 68.87s/it]

Epoch: 7 Training Loss: 3.535999 Validation Loss: 4.186358

80%| | 8/10 [09:11<02:17, 68.85s/it]

Epoch: 8            Training Loss: 3.339648            Validation Loss: 4.222634

90%| | 9/10 [10:19<01:08, 68.71s/it]

Epoch: 9            Training Loss: 3.161209            Validation Loss: 4.240358

100%|| 10/10 [11:28<00:00, 68.65s/it]

Epoch: 10           Training Loss: 2.970675           Validation Loss: 4.220605

Test Loss: 4.089101

Test Accuracy: 8% (74/836)

In [14]: evaluate\_model(resaug\_loaders, model\_scratch, optimizer\_scratch, criterion\_scratch, use

10%|            | 1/10 [01:32<13:56, 92.92s/it]

Epoch: 1            Training Loss: 3.991425            Validation Loss: 4.126298

Validation loss decreased. Saving model ...

20%|            | 2/10 [03:05<12:21, 92.68s/it]

Epoch: 2            Training Loss: 3.851191            Validation Loss: 3.978856

Validation loss decreased. Saving model ...

30%|            | 3/10 [04:37<10:48, 92.59s/it]

Epoch: 3            Training Loss: 3.747243            Validation Loss: 3.893922

Validation loss decreased. Saving model ...

40%|            | 4/10 [06:09<09:14, 92.44s/it]

Epoch: 4            Training Loss: 3.656609            Validation Loss: 3.948472

50%|            | 5/10 [07:42<07:43, 92.64s/it]

Epoch: 5            Training Loss: 3.579864            Validation Loss: 3.859303

Validation loss decreased. Saving model ...

60%| | 6/10 [09:15<06:11, 92.82s/it]

|   |                         |                           |
|---|-------------------------|---------------------------|
| Epoch: 6                                    | Training Loss: 3.501025 | Validation Loss: 3.808497 |
| Validation loss decreased. Saving model ... |                         |                           |

70%| | 7/10 [10:48<04:38, 92.91s/it]

|   |                         |                           |
|---|-------------------------|---------------------------|
| Epoch: 7                                    | Training Loss: 3.428164 | Validation Loss: 3.761371 |
| Validation loss decreased. Saving model ... |                         |                           |

80%| | 8/10 [12:22<03:05, 92.94s/it]

|   |                         |                           |
|---|-------------------------|---------------------------|
| Epoch: 8                                    | Training Loss: 3.323892 | Validation Loss: 3.744287 |
| Validation loss decreased. Saving model ... |                         |                           |

90%| | 9/10 [13:55<01:33, 93.09s/it]

|          |                         |                           |
|----------|-------------------------|---------------------------|
| Epoch: 9 | Training Loss: 3.254543 | Validation Loss: 3.752984 |
|----------|-------------------------|---------------------------|

100%|| 10/10 [15:29<00:00, 93.30s/it]

|   |                         |                           |
|---|-------------------------|---------------------------|
| Epoch: 10                                   | Training Loss: 3.149271 | Validation Loss: 3.715734 |
| Validation loss decreased. Saving model ... |                         |                           |

Test Loss: 3.743450

Test Accuracy: 13% (113/836)

In [34]: evaluate\_model(craug\_loaders, model\_scratch, optimizer\_scratch, criterion\_scratch, use\_

|   |                         |                           |
|---|-------------------------|---------------------------|
| Epoch: 1                                    | Training Loss: 3.981338 | Validation Loss: 4.000509 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 2                                    | Training Loss: 3.779310 | Validation Loss: 3.965888 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 3                                    | Training Loss: 3.660987 | Validation Loss: 3.906343 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 4                                    | Training Loss: 3.547598 | Validation Loss: 3.859987 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 5                                    | Training Loss: 3.476992 | Validation Loss: 3.784964 |
| Validation loss decreased. Saving model ... |                         |                           |



|   |                         |                           |
|---|-------------------------|---------------------------|
| Epoch: 6                                    | Training Loss: 3.411273 | Validation Loss: 3.839303 |
| Epoch: 7                                    | Training Loss: 3.321054 | Validation Loss: 4.000874 |
| Epoch: 8                                    | Training Loss: 3.236289 | Validation Loss: 3.849416 |
| Epoch: 9                                    | Training Loss: 3.131379 | Validation Loss: 3.766074 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 10                                   | Training Loss: 3.050062 | Validation Loss: 3.773075 |
| Test Loss: 3.715546                         |                         |                           |

Test Accuracy: 15% (131/836)

---

#### ## Step 4: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

##### 1.1.12 (IMPLEMENTATION) Specify Data Loaders for the Dog Dataset

Use the code cell below to write three separate [data loaders](#) for the training, validation, and test datasets of dog images (located at dogImages/train, dogImages/valid, and dogImages/test, respectively).

If you like, **you are welcome to use the same data loaders from the previous step**, when you created a CNN from scratch.

```
In [2]: ## TODO: Specify data loaders
import numpy as np
import torch
import torchvision.models as models
import os
from torchvision import datasets, transforms
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

### TODO: Write data loaders for training, validation, and test sets
## Specify appropriate transforms, and batch_sizes
def load_transfer_data_in_DataLoaders(data_dir, tr_data_transform = transforms.ToTensor():
    # Assemble dir paths
    net_stage = ['train', 'valid', 'test']
    # Load Images dataset using ImageFolder
    data_transfer = { item: datasets.ImageFolder(os.path.join(data_dir, item), transform
        for item in net_stage}

    # Print out info
    for stage in net_stage:
        print('Num images in',stage, ':', len(data_transfer[stage]) )

    # prepare data loaders
```

```

        loaders_transfer = {item : torch.utils.data.DataLoader(data_transfer[item], batch_size=batch_size,
                                                                num_workers=num_workers, shuffle = True) for item in data_transfer.keys()}

    return loaders_transfer

In [3]: def create_transfer_transforms(img_size):
        # Define transforms on the data
        tr_data_transform = transforms.Compose([
            transforms.Resize((img_size,img_size)),
            transforms.RandomHorizontalFlip(0.05),
            transforms.RandomRotation(30),
            transforms.RandomAffine(10),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])
        ])
        tr_crop_data_transform = transforms.Compose([
            transforms.RandomCrop(img_size),
            transforms.RandomHorizontalFlip(0.05),
            transforms.RandomRotation(30),
            transforms.RandomAffine(10),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])
        ])
        vt_data_transform = transforms.Compose([
            transforms.Resize((img_size,img_size)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])
        ])
        vt_crop_data_transform = transforms.Compose([
            transforms.CenterCrop(img_size),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])
        ])
        return tr_data_transform, tr_crop_data_transform, vt_data_transform, vt_crop_data_transform

In [4]: img_size = 224
        tr_transfer_data_transform, tr_crop_transfer_data_transform, vt_transfer_data_transform,
        transfer_res_loaders = load_transfer_data_in_DataLoaders('/data/dog_images', vt_transfer_data_transform)
        transfer_cr_loaders = load_transfer_data_in_DataLoaders('/data/dog_images', vt_crop_transfer_data_transform)

        transfer_resaug_loaders = load_transfer_data_in_DataLoaders('/data/dog_images', tr_transfer_data_transform)
        transfer_craug_loaders = load_transfer_data_in_DataLoaders('/data/dog_images', tr_crop_transfer_data_transform)

Num images in train : 6680
Num images in valid : 835

```

```

Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836

```

```

In [5]: img_size = 448
        tr_transfer_data_transform_448, tr_crop_transfer_data_transform_448, vt_transfer_data_tr

        res_loaders_448 = load_transfer_data_in_DataLoaders('/data/dog_images', vt_transfer_data
        cr_loaders_448 = load_transfer_data_in_DataLoaders('/data/dog_images', vt_crop_transfer

        resaug_loaders_448 = load_transfer_data_in_DataLoaders('/data/dog_images', tr_transfer_
        craug_loaders_448 = load_transfer_data_in_DataLoaders('/data/dog_images', tr_crop_transf

Num images in train : 6680
Num images in valid : 835
Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836
Num images in train : 6680
Num images in valid : 835
Num images in test : 836

```

### 1.1.13 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify dog breed. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [6]: import torchvision.models as models
        import torch.nn as nn

        # check if CUDA is available
        use_cuda = torch.cuda.is_available()

        ## TODO: Specify model architecture
        VGG16 = models.vgg16(pretrained=True)
        IncV3 = models.inception_v3(pretrained=True)

```

```
RNET50 = models.resnet50(pretrained=True)
models = {'VGG16': VGG16, 'IncV3': IncV3, 'RNET50': RNET50}
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.torch/models/vgg16
100%|| 553433881/553433881 [00:07<00:00, 73288210.20it/s]
Downloading: "https://download.pytorch.org/models/inception_v3_google-1a9a5a14.pth" to /root/.torch/models/inception_v3
100%|| 108857766/108857766 [00:01<00:00, 72424541.94it/s]
Downloading: "https://download.pytorch.org/models/resnet50-19c8e357.pth" to /root/.torch/models/resnet50
100%|| 102502400/102502400 [00:01<00:00, 86107776.63it/s]
```

```
In [6]: model_transfer = VGG16
        if use_cuda:
            model_transfer = model_transfer.cuda()
```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:** Three architectures are evaluated: vgg16, inception v3 and resnet. Vgg16 trades-off accuracy with ease of computation. The 3x3 kernel allows increasing the depth of the CNN convs providing better detection than previous network. Thus, it is a good choice starting with it and moving into other architectures when higher accuracy is required. Resnet is computationally complex but not too demanding to train. The architecture has delivered better results when assessed in dog/human detector at step 2. I have used the nets freezing the feature / conv layers, and changed only the linear output to adapt the fully connected layer to fit the classes at hand.

#### 1.1.14 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and the optimizer as `optimizer_transfer` below.

```
In [7]: import torch.optim as optim

        ### TODO: select loss function
        criterion_transfer = nn.CrossEntropyLoss()

        ### TODO: select optimizer
        optimizer_transfer = optim.SGD(model_transfer.parameters(), lr=0.01)
```

#### 1.1.15 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath `'model_transfer.pt'`.

```
In [11]: def freeze_feature_layers(model_transfer, output_classes, use_cuda):
        for param in model_transfer.features.parameters():
            param.requires_grad = False

        # modify last classifier layer to fit our output
```

```

        model_transfer.classifier[-1] = nn.Linear(model_transfer.classifier[-1].in_features

        # move to gpu is available
        if use_cuda:
            model_transfer = model_transfer.cuda()
        return model_transfer

In [12]: # train the model
dog_breeds = 133
model_transfer = VGG16
if use_cuda:
    model_transfer = model_transfer.cuda()

model_transfer = freeze_feature_layers(model_transfer, dog_breeds, use_cuda)
model_transfer = train(10, transfer_resaug_loaders, model_transfer, optimizer_transfer,
                        use_cuda, 'model_transfer.pt')

# load the model that got the best validation accuracy (uncomment the line below)
model_transfer.load_state_dict(torch.load('model_transfer.pt'))

Epoch: 1      Training Loss: 3.992542      Validation Loss: 2.379578
Validation loss decreased. Saving model ...
Epoch: 2      Training Loss: 2.157105      Validation Loss: 1.429146
Validation loss decreased. Saving model ...
Epoch: 3      Training Loss: 1.570267      Validation Loss: 1.162556
Validation loss decreased. Saving model ...
Epoch: 4      Training Loss: 1.301861      Validation Loss: 1.085887
Validation loss decreased. Saving model ...
Epoch: 5      Training Loss: 1.109071      Validation Loss: 0.924187
Validation loss decreased. Saving model ...
Epoch: 6      Training Loss: 0.998924      Validation Loss: 0.883514
Validation loss decreased. Saving model ...
Epoch: 7      Training Loss: 0.898183      Validation Loss: 0.905753
Epoch: 8      Training Loss: 0.821752      Validation Loss: 0.868860
Validation loss decreased. Saving model ...
Epoch: 9      Training Loss: 0.739293      Validation Loss: 0.848797
Validation loss decreased. Saving model ...
Epoch: 10     Training Loss: 0.696382      Validation Loss: 0.824698
Validation loss decreased. Saving model ...

In [10]: import torch.optim as optim
model_transfer_RNET = RNET50

### TODO: select loss function
criterion_transfer = nn.CrossEntropyLoss()

### TODO: select optimizer

```

```

optimizer_transfer = optim.SGD(model_transfer_RNET.parameters(), lr=0.01)

# train the model
def freeze_RNET_layers(model_transfer, output_classes, use_cuda):
    for param in model_transfer.parameters():
        param.requires_grad = True

    # modify last classifier layer to fit our output
    model_transfer.classifier = nn.Linear(2048, output_classes)

    # move to gpu is available
    if use_cuda:
        model_transfer = model_transfer.cuda()
    return model_transfer

dog_breeds = 133
if use_cuda:
    model_transfer_RNET = model_transfer_RNET.cuda()

model_transfer_RNET = freeze_RNET_layers(model_transfer_RNET, dog_breeds, use_cuda)
model_transfer_RNET = train(10, transfer_resaug_loaders, model_transfer_RNET, optimizer_transfer,
                             use_cuda, 'model_transfer_RNET.pt')

# load the model that got the best validation accuracy (uncomment the line below)
model_transfer_RNET.load_state_dict(torch.load('model_transfer_RNET.pt'))

```

|   |                         |                           |
|---|-------------------------|---------------------------|
| Epoch: 1                                    | Training Loss: 3.320516 | Validation Loss: 1.316821 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 2                                    | Training Loss: 1.074966 | Validation Loss: 0.979682 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 3                                    | Training Loss: 0.676241 | Validation Loss: 0.751343 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 4                                    | Training Loss: 0.463238 | Validation Loss: 0.714227 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 5                                    | Training Loss: 0.345936 | Validation Loss: 0.705079 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 6                                    | Training Loss: 0.256075 | Validation Loss: 0.718243 |
| Epoch: 7                                    | Training Loss: 0.198520 | Validation Loss: 0.686255 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 8                                    | Training Loss: 0.176412 | Validation Loss: 0.678906 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 9                                    | Training Loss: 0.129701 | Validation Loss: 0.678232 |
| Validation loss decreased. Saving model ... |                         |                           |
| Epoch: 10                                   | Training Loss: 0.109922 | Validation Loss: 0.687057 |

### 1.1.16 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [12]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    model.eval()
    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))
```

```
In [17]: test(transfer_resaug_loaders, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 0.866053

Test Accuracy: 74% (620/836)

```
In [14]: test(transfer_resaug_loaders, model_transfer_RNET, criterion_transfer, use_cuda)
```

Test Loss: 0.645941

Test Accuracy: 82% (693/836)



Sample Human Output

### 1.1.17 (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan hound, etc) that is predicted by your model.

```
In [7]: ### TODO: Write a function that takes a path to an image as input
### and returns the dog breed that is predicted by the model.
data_transfer = transfer_res_loaders
# list of class names by index, i.e. a name can be accessed like class_names[0]
class_names = [item[4:].replace("_", " ") for item in data_transfer['train'].dataset.class_names]

def predict_breed_transfer(img_path, transforms, model_transfer, class_names, use_cuda):
    # load the image and return the predicted breed
    image = Image.open(img_path).convert('RGB')
    img = transforms(image)
    img = img.unsqueeze(0)
    if use_cuda:
        img = img.cuda()

    model_transfer.eval()
    output = model_transfer(img)
    _, pred = torch.max(output, 1)
    return class_names[pred]
```

---

#### ## Step 5: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then, - if a **dog** is detected in the image, return the predicted breed. - if a **human** is detected in the image, return the resembling dog breed. - if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `human_detector` functions developed above. You are **required** to use your CNN from Step 4 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!



### 1.1.18 (IMPLEMENTATION) Write your Algorithm

```
In [8]: def dog_detector(img_path, img_size, model = VGG16):  
        ## TODO: Complete the function.  
        output = VGG16_predict(img_path, img_size, model)  
        return output in range(151, 268)
```

```
In [22]: ### TODO: Write your algorithm.  
        ### Feel free to use as many code cells as needed.
```

```
def run_app(img_path):  
    ## handle cases for a human face, dog, and neither  
    img_size = 224  
    img = Image.open(img_path)  
    plt.figure()  
    plt.title(img_path)  
    plt.imshow(img)  
    plt.show()  
    vt_data_transform = transforms.Compose([  
        transforms.Resize((img_size, img_size)),  
        transforms.ToTensor(),  
        transforms.Normalize(mean=[0.485, 0.456, 0.406],  
                              std=[0.229, 0.224, 0.225])  
    ])  
    if dog_detector(img_path, img_size, VGG16):  
        print('Dog detected. Breed', predict_breed_transfer(img_path, vt_data_transform, class_names, use))  
    elif face_detector(img_path):  
        print('Human detected, looks like a : ', predict_breed_transfer(img_path, vt_data_transform, class_names, use))  
    else:  
        print('Bad detection. What is that?')
```

---

#### ## Step 6: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that *you* look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

### 1.1.19 (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

**Question 6:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

**Answer:** (Three possible points for improvement):

Three points for improvement are: - Use face detector to mask and isolate faces so background has less impact - Determine or improve face detectors to determine whether there are more faces

in the images. Same for dogs - Get more images to improve training dataset, especially of dogs in sideview or not centered - Train for more epochs, use adaptive learning rates - Use batchnormalization layers - Try different networks that admit higher resolution - Define functions to mask the silhouette of the dog / human face so the network always focus of face + body features

```
In [10]: import numpy as np
         from glob import glob

         # load filenames for human and dog images
         human_files = np.array(glob("/data/lfw/*/"))
         dog_files = np.array(glob("/data/dog_images/*/"))

         # print number of images in each dataset
         print('There are %d total human images.' % len(human_files))
         print('There are %d total dog images.' % len(dog_files))
```

There are 13233 total human images.  
There are 8351 total dog images.

```
In [12]: import torch.optim as optim
         model_transfer_RNET = RNET50

         ### TODO: select loss function
         criterion_transfer = nn.CrossEntropyLoss()

         ### TODO: select optimizer
         optimizer_transfer = optim.SGD(model_transfer_RNET.parameters(), lr=0.01)

         # train the model
         def freeze_RNET_layers(model_transfer, output_classes, use_cuda):
             for param in model_transfer.parameters():
                 param.requires_grad = True

             # modify last classifier layer to fit our output
             model_transfer.classifier = nn.Linear(2048, output_classes)

             # move to gpu is available
             if use_cuda:
                 model_transfer = model_transfer.cuda()
             return model_transfer


         dog_breeds = 133
         if use_cuda:
             model_transfer_RNET = model_transfer_RNET.cuda()

         model_transfer_RNET = freeze_RNET_layers(model_transfer_RNET, dog_breeds, use_cuda)
```

```
# load the model that got the best validation accuracy (uncomment the line below)
model_transfer_RNET.load_state_dict(torch.load('model_transfer.pt'))
```


```
In [23]: ## TODO: Execute your algorithm from Step 6 on
## at least 6 images on your computer.
## Feel free to use as many code cells as needed.
from PIL import Image
import matplotlib.pyplot as plt
%matplotlib inline

## suggested code, below
for file in np.hstack((human_files[:5], dog_files[:5])):
    run_app(file)
```




Human detected, looks like a : Dogue de bordeaux

/data/lfw/Alex\_Corretja/Alex\_Corretja\_0001.jpg

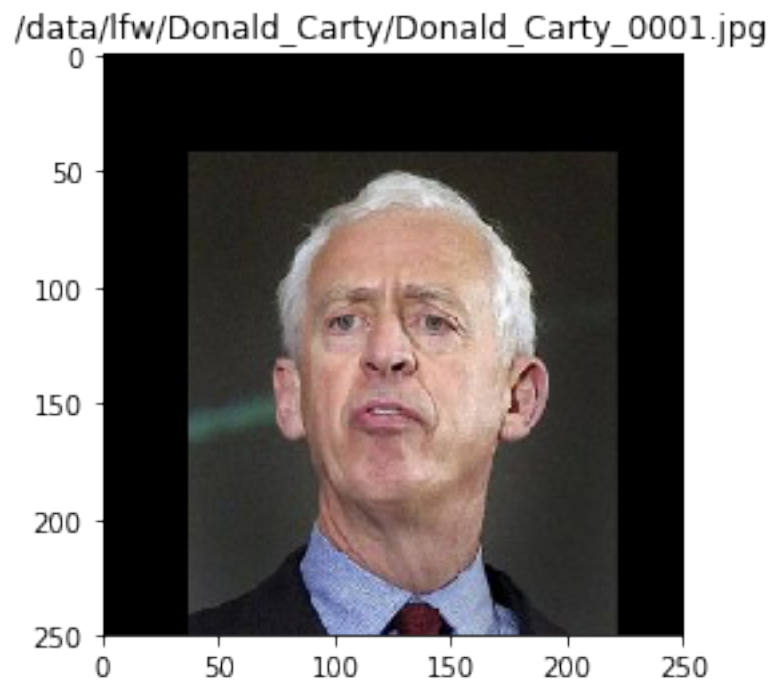


Human detected, looks like a : Lakeland terrier


/data/lfw/Daniele\_Bergamin/Daniele\_Bergamin\_0001.jpg




Human detected, looks like a : American water spaniel



Human detected, looks like a : Dogue de bordeaux




Human detected, looks like a : Dogue de bordeaux




Bad detection. What is that?

/data/dog\_images/train/103.Mastiff/Mastiff\_06826.jpg



Dog detected. Breed Mastiff

/data/dog\_images/train/103.Mastiff/Mastiff\_06871.jpg






Dog detected. Breed Mastiff



Bad detection. What is that?






Dog detected. Breed Mastiff


```
In [25]: ## TODO: Execute your algorithm from Step 6 on
        ## at least 6 images on your computer.
        ## Feel free to use as many code cells as needed.
        from PIL import Image
        import matplotlib.pyplot as plt
        %matplotlib inline

        ## suggested code, below
        for file in np.hstack((human_files[50], dog_files[50])):
            run_app(file)
```

/data/lfw/Julianne\_Moore/Julianne\_Moore\_0003.jpg



Human detected, looks like a : Dogue de bordeaux

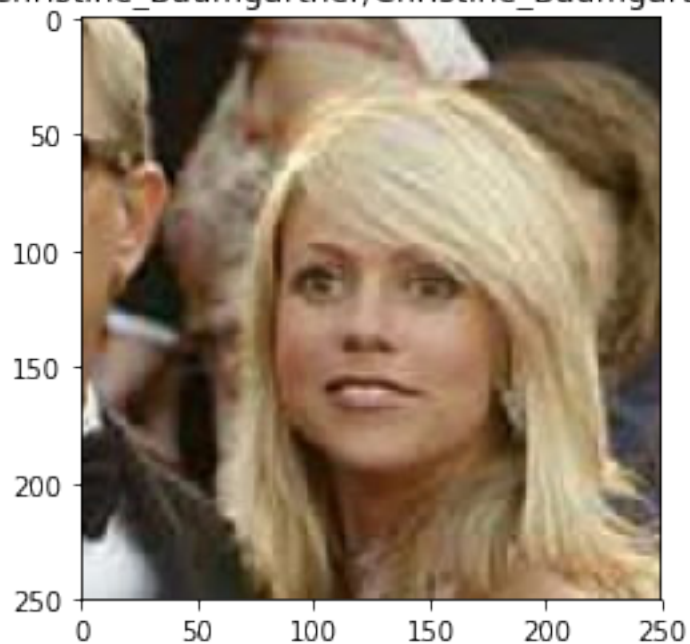


Dog detected. Breed Mastiff

```
In [26]: ## TODO: Execute your algorithm from Step 6 on
        ## at least 6 images on your computer.
        ## Feel free to use as many code cells as needed.
        from PIL import Image
        import matplotlib.pyplot as plt
        %matplotlib inline


        ## suggested code, below
        for file in np.hstack((human_files[150], dog_files[150])):
            run_app(file)
```

/data/lfw/Christine\_Baumgartner/Christine\_Baumgartner\_0002.jpg



Human detected, looks like a : American water spaniel


/data/dog\_images/train/055.Curly-coated\_retriever/Curly-coated\_retriever\_03873.jpg



Dog detected. Breed Curly-coated retriever

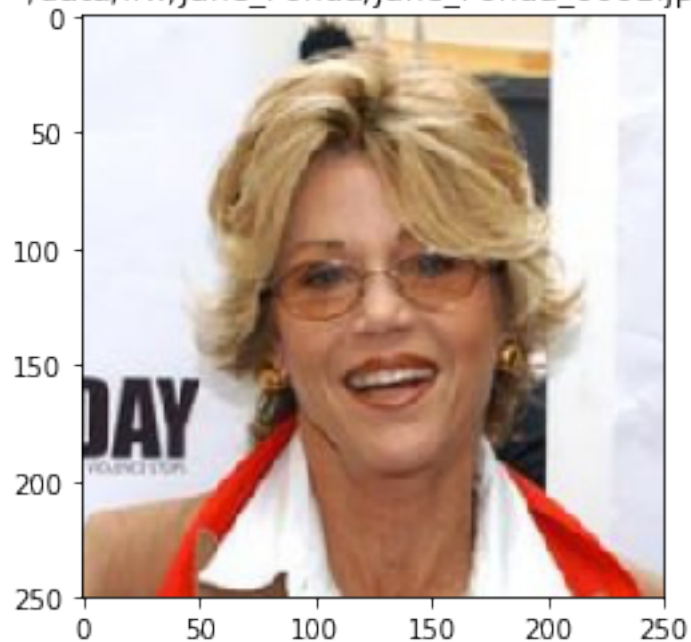
```
In [27]: ## TODO: Execute your algorithm from Step 6 on
        ## at least 6 images on your computer.
        ## Feel free to use as many code cells as needed.
        from PIL import Image
        import matplotlib.pyplot as plt
        %matplotlib inline

        ## suggested code, below
        for file in np.hstack((human_files[250:252], dog_files[250:252])):
            run_app(file)
```




Bad detection. What is that?

/data/lfw/Jane\_Fonda/Jane\_Fonda\_0001.jpg



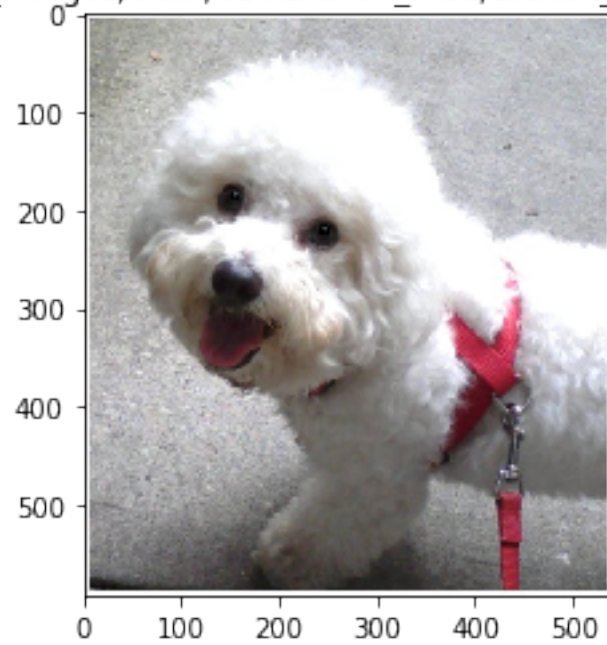
Human detected, looks like a : Afghan hound

/data/dog\_images/train/024.Bichon\_frise/Bichon\_frise\_01743.jpg



Dog detected. Breed Bichon frise

/data/dog\_images/train/024.Bichon\_frise/Bichon\_frise\_01762.jpg



Dog detected. Breed Bichon frise

In [ ]: